


The Sleeping Barber problem is attributed to Dijkstra, and is based on the following scenario:

A barbershop contains a barber, the barber's chair, and N chairs for waiting customers. When there are no customers, the barber sits in his chair and sleeps. When a customer arrives, (a) if the barber is sleeping, the customer awakens the barber and sits down for a haircut; (b) if the barber is busy and there is a free chair, the customer sits down and waits; (c) if the barber is busy and there are no free chairs, the customer leaves immediately.



More specifically, there are 10 customers and 1 barber, each represented by a thread, and a monitor representing the barbershop with two methods, `barber()` and `customer()`.

- The barber thread calls `barber()`, which never returns. (i.e. the barber never leaves the shop) The barber sleeps until woken by a customer, and then gives haircuts (average time 1.2 seconds, exponentially distributed) until there are no customers left, after which he goes to sleep again.
- Each customer thread loops, alternately sleeping for a period of time (average 10 seconds, exponentially distributed) and then calling `customer()`, which returns after the customer leaves the shop. (immediately if all chairs are full, or after getting a haircut)

Question 1 – Synchronization

Part 1: Provide pseudo-code for a monitor which has two methods, `barber()` and `customer()`, modeling a barbershop with 4 waiting chairs and one barber's chair.

Remember the following characteristics of the monitor definition used in this class:

- Only one thread can be in the monitor at a time; threads enter the monitor at the beginning of a method or when returning from `wait()`, and leave the monitor by returning from a method or entering `wait()`.
- This means there is no preemption – a thread in a method executes without interruption until it returns or waits.
- When thread A calls `signal()` to release thread B from waiting on a condition variable, you don't know whether B will run before or after some thread C that tries to enter the method at the same time. (that's why they're called *race* conditions)
- You don't need a separate mutex – this is a monitor, not a pthreads translation of a monitor.

Part 2: Illustrate the operation of this monitor using the graphical notation introduced in class.

Further submission instructions: your pseudo-code in Part 1 should be specific, rather than just giving vague generalities. In particular you should give enough information so that a reader can tell whether your code implements the same algorithm as your pseudo-code. (e.g. at a bare minimum you have to identify the condition variables and when they are waited on and when they are signaled)

You may submit Part 1 as a text file (`q1-part1.txt`) or as part of a PDF file including part 2; in either case the file should be added to your repository and committed. Part 2 should be submitted as a PDF file; it doesn't matter whether it is hand-drawn and scanned (as long as I can read it) or prepared nicely with something like OpenOffice Draw. Drawings on paper, legibly labeled with your name, may be slipped under my door as well.

Question 2 – POSIX Threads

There is a straightforward translation from monitor pseudo-code to POSIX thread primitives:

1. Create a per-object mutex, m , of type `pthread_mutex_t` which is locked on entry to each method and unlocked on exit. (be careful when using multiple exits) Actually, since there's only one object – the barbershop – there should only be one mutex.
2. Condition variables translate directly to objects of type `pthread_cond_t`; `C.signal()` and `C.broadcast()` become `pthread_cond_signal(C)` and `pthread_cond_broadcast(C)`;
3. The monitor mutex must be passed to wait calls; thus `C.wait()` becomes `pthread_cond_wait(C,m)`.

For this exercise we will create a singleton monitor, using global variables instead of object variables, and functions rather than object methods.

You will use a single file, *homework.c*, for both question 2 and 3, using conditional compilation to separate the code for the two; code for question 2 will be compiled with the script *compile-q2.sh* creating an executable named *homework-q2*. The startup code for this question will go in the function *q2()*, and will do the following:

- Initialize the monitor objects. Note that mutexes and condition variables may be initialized either statically or dynamically:

```
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER; /* static init */
pthread_cond_t C = PTHREAD_COND_INITIALIZER;

pthread_mutex_t m;                               /* dynamic */
pthread_mutex_init(&m, NULL);                     /* NULL = default params */

pthread_cond_t C;
pthread_cond_init(&C, NULL);
```

- Create $N=10$ customer threads; each thread will loop doing the following:
... sleep for `random(T seconds)`...
`customer()`

You are provided with a sleep function, *sleep_exp(T)*, where T is a floating point number giving the mean sleep time in seconds. Each thread will need to know its thread number, from 1 to 10; there is a comment in the code describing how to pass this value when starting a thread.

- Call *wait_until_done()*, which will sleep until a command-line-provided timeout or until the user types ^C.

Running question 2: The *homework-q2* command is used as follows:

```
./homework-q2 [-speedup #f] [#t]
```

where *#t* is a total number of seconds the homework should run for, and *#f* is a speedup factor. (e.g. *./homework-q2 100* would run for 100 seconds; *./homework-q2 -speedup 2.5 100* would do the same amount of work, but run 2.5 times faster, completing 100 simulated seconds in 40 real seconds.)

To use the debug script provided for this question, you will need to print the following lines as your code executes:

```
DEBUG: TTT customer # enters shop
DEBUG: TTT customer # starts haircut
DEBUG: TTT customer # leaves shop
DEBUG: TTT barber wakes up
DEBUG: TTT barber goes to sleep
```

where “TTT” is a floating point timestamp returned by the *timestamp()* function.

If you redirect the output of the command into a file, you should be able to use the *verify-q2* script to determine whether or not your implementation obeyed all the rules:

```
./homework-q2 ... > q2.out
./verify-q2 q2.out
```

Note that if you run your program for a short period of time, it is less likely to break any rules even if it is incorrect.

I suggest running Q2 for a while with a high speedup, and possibly performing some work in another terminal window on the same machine (to disturb the thread scheduling order) in order to determine whether it operates correctly and whether it deadlocks.

Question 3 – Discrete Event Simulation

For this question you will compile your code using the *compile-q3* script, which will build it with a framework (from *misc.c*, using the GNU Pth library) for discrete-event simulation. What this means is that your code will run in *simulated time* – basically the thread library “skips” time forward whenever threads are sleeping, and stops the clock when a thread is running. For simulations of small, slow systems (like ours) this results in a simulation running much faster than real time; for fast, complex systems (e.g. simulating operation of an integrated circuit) the simulation might run thousands of times slower than real time.

The simulation library is designed so that it is compatible with Pthreads operations, so you should be able to compile and run the same monitor code you used in Question 2. In addition several functions have been provided for gathering statistics:

```
void *counter = stat_counter();
stat_count_incr(counter);
stat_count_decr(counter);
double val = stat_count_mean(counter);

void *timer = stat_timer();
stat_timer_start(timer);
stat_timer_stop(timer);
double val = stat_timer_mean(timer)
```

A counter tracks an integer variable (e.g. the number of customers waiting in the shop) and provides its average value over time. A timer tracks the interval between a single thread calling *start()* and *stop()*, and provides the average value of these measured intervals.

Run your code for at least 1000 seconds (10000 would be better) of simulated time and measure average values for the following:

- fraction of customer visits result in turning away due to a full shop (you'll have to keep your own counter for this one)
- average time spent in the shop (including haircut) by a customer who does not find a full shop
- average number of customers in the shop (including the barber's chair)
- fraction of time someone is sitting in the barber's chair (hint – use a *stat_counter* with value 0 for empty and 1 for full)